

LAB 4

Logistic Regression for Classification

PART A: Prerequisite for Linear Regression implementation

1. Plot the attached dataset data1.csv using scatter plot. There is a target feature with discrete values 0,1. If the target feature is 1, the samples should be shown as red circle. If the target feature is 0, the samples should be shown as green x.

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from scipy.optimize import fmin_tnc
import scipy.optimize as op

from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score,
classification_report, roc_auc_score, roc_curve
plt.style.use("Solarize_Light2")
import random
np.random.seed(9)
np.set_printoptions(suppress=True)
import math
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df_data1 = pd.read_csv("data1.csv")
df_data1
```


```
Out[2]:
```

	x1	x2	y
0	4.5192	2.6487	1
1	2.4443	1.5438	1
2	4.2409	1.8990	1
3	5.8097	2.4711	1
4	6.4423	3.3590	1
...

	x1	x2	y
95	5.9868	7.3641	0
96	4.6711	6.2592	0
97	7.5810	8.3703	0
98	4.6457	8.5676	0
99	4.6457	8.1676	0

```
In [3]: x1 = df_data1['x1']
x2 = df_data1['x2']
y = df_data1[['y']]
col = np.where(y == 1, 'r', 'b')
col = [i[0] for i in col]
```

```
In [4]: plt.scatter(x1, x2, c=col, s=10, linewidth=3)
plt.xlabel("x1")
plt.ylabel("x2")
plt.show()
```



2. Plot a line $y=(2x+3)$


```

In [5]: y1 = [(2*i+3) for i in list(x1)]
y2 = [(2*i+3) for i in list(x2)]
fig = plt.figure()
fig.set_figheight(5)
fig.set_figwidth(20)
plt.subplot(1, 2, 1)
plt.plot(x1, y1, color='green', label = "line y1 = 2(x1)+3", marker='o',
markerfacecolor='black', markersize=3)
plt.xlabel('x - axis')
plt.ylabel('y - axis')
plt.title('line y = 2(x1)+3')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(x2, y2, color='blue', label = "line y2 = 2(x2)+3", marker='o',
markerfacecolor='black', markersize=3)
plt.xlabel('x - axis')
plt.ylabel('y - axis')
plt.title('line y = 2(x2)+3')
plt.legend()

```

Out[5]: <matplotlib.legend.Legend at 0x248af67bf10>



3. Define a function `sigmoid(z)` that takes one parameter `z` and computes $1/(1+e^{-z})$. Create a vector `V` with 10 values randomly in the range `[-1000,1000]`. Transform `V` to `V'` that consists of respective sigmoid values using the defined function. Observe the range of output values in `V'`.

```

In [6]: def sigmoid(z): # Activation function used to map any real value between 0
and 1
return (1 / (1 + np.exp(-z)))

```

```

In [7]: v = [round(random.uniform(-1000, 1000)) for i in range(10)]
v_ = [sigmoid(z) for z in v]
print(v)
print(v_)

```

```

[-312, -543, -75, 740, 783, 239, 436, 463, -373, 987]

```

4. Define a function hypothesis(theta, X) that takes two vectors as parameters, theta and X. If $\text{sigmoid}(\text{theta} \cdot X) \geq 0.5$, output 0 else output 1.

5. Define a function `cost(theta,X,y)` to compute the error
 $\text{Error} = \frac{1}{m} \sum -y_i \log(h_{\theta}(x_i)) - (1-y_i) \log(1-h_{\theta}(x_i))$
 Where x_i is the i th sample and y_i is the i th label, $h_{\theta}(x_i)$ is the hypothesis(θ, x_i)

PART B : Implementation of logistic regression

- 6. Implement gradient descent algorithm for logistic regression in data set `loan_data.csv`.
 - read training dataset
 - convert dataset to a feature matrix X
 - normalizing feature matrix X
 - stack columns with all ones in feature matrix
 - target feature to be taken in a separate vector
 - initial theta values
 - gradient descent updation
 - Display estimated theta values and number of iterations to converge
 - Predict for test data

read training dataset

Out[10]:	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	L
	0	LP001002	Male	No	0	Graduate	No	5849	0.0

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	Loan_Amount_Term
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	36
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	36
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	36
4	LP001008	Male	No	0	Graduate	No	6000	0.0	36
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	36
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	36
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	36
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	36
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	36

In [11]: `loan_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                614 non-null    object
1   Gender                 601 non-null    object
2   Married                611 non-null    object
3   Dependents             599 non-null    object
4   Education              614 non-null    object
5   Self_Employed          582 non-null    object
6   ApplicantIncome        614 non-null    int64
7   CoapplicantIncome      614 non-null    float64
8   LoanAmount             592 non-null    float64
9   Loan_Amount_Term       600 non-null    float64
10  Credit_History         564 non-null    float64
11  Property_Area          614 non-null    object
12  Loan_Status            614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```


In [12]: `print("Column Name\tNull values")`
`print(loan_data.isnull().sum())`

```
Column Name      Null values
Loan_ID          0
Gender           13
Married          3
Dependents       15
Education        0
Self_Employed    32
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       22
Loan_Amount_Term 14
Credit_History  50
Property_Area    0
Loan_Status      0
dtype: int64
```

Filling up null values

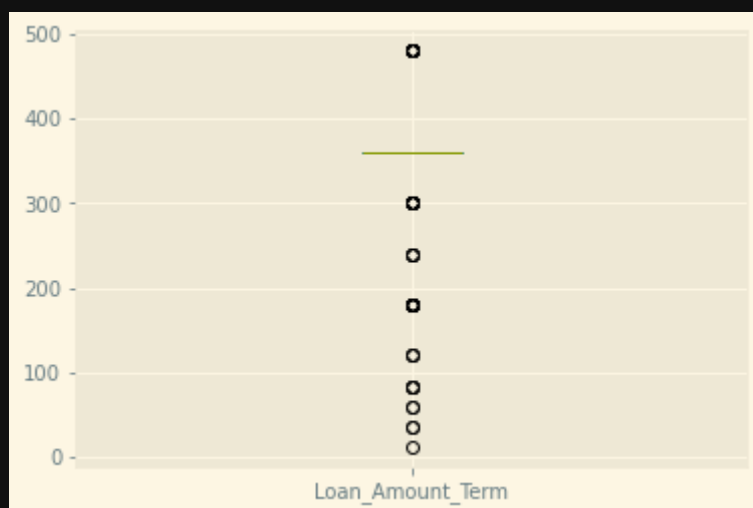
In [13]: `loan_data["LoanAmount"].plot(kind = 'box') #as LoanAmount is skewed, we will fill it with median`

Out[13]: <AxesSubplot:>




```
In [14]: loan_data["Loan_Amount_Term"].plot(kind = 'box') #as Loan_Amount_Term is skewed, we will fill it with median
```

Out[14]: <AxesSubplot:>



```
In [15]: loan_data["Dependents"].replace("3+", "3" , inplace = True)
loan_data["Dependents"] = pd.to_numeric(loan_data["Dependents"],
errors='coerce')
loan_data["Dependents"].plot(kind = 'box') #as Dependents is skewed, we will fill it with median
```

Out[15]: <AxesSubplot:>



```
In [16]: loan_data["Dependents"].unique()
```

```
Out[16]: array([ 0.,  1.,  2.,  3., nan])
```

```
In [17]: loan_data["LoanAmount"].fillna(loan_data['LoanAmount'].median(), inplace =
True)
loan_data["Loan_Amount_Term"].fillna(loan_data['LoanAmount'].median(),
inplace = True)
loan_data["Dependents"].fillna(loan_data['Dependents'].median(), inplace =
True)
loan_data["Dependents"] = loan_data["Dependents"].astype('Int64')
```

```
In [18]: #Other columns are categorical data, so we will fill up null values with
mode
loan_data["Gender"].fillna("No Gender", inplace = True)
loan_data["Married"].fillna(loan_data['Married'].mode()[0], inplace = True)
loan_data["Dependents"].fillna(loan_data['Dependents'].mode()[0], inplace =
True)
loan_data["Self_Employed"].fillna(loan_data['Self_Employed'].mode()[0],
inplace = True)
loan_data["Credit_History"].fillna(loan_data['Credit_History'].mode()[0],
inplace = True)
```

```
In [19]: print("Column Name\tNull values")
print(loan_data.isnull().sum())
```

Column Name	Null values
Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0
Property_Area	0
Loan_Status	0

dtype: int64

```
In [20]: categorical_columns = ["Gender", "Married", "Education", "Self_Employed",
"Credit_History", "Property_Area"]
y_categorical_columns = ["Loan_Status"]
numerical_columns = ["Dependents", "LoanAmount", "Loan_Amount_Term",
"ApplicantIncome", "CoapplicantIncome"]
for col in categorical_columns:
    print(col, "\t\t", list(loan_data[col].unique()))
```

Gender	['Male', 'Female', 'No Gender']
Married	['No', 'Yes']
Education	['Graduate', 'Not Graduate']

```

Self_Employed      ['No', 'Yes']
Credit_History      [1.0, 0.0]
Property_Area      ['Urban', 'Rural', 'Semiurban']

```

```

In [21]: loan_data["Loan_Amount_Term"] =
loan_data["Loan_Amount_Term"].astype('Int64')
loan_data["Credit_History"] = loan_data["Credit_History"].astype('str')
loan_data["Loan_Status"] = loan_data["Loan_Status"].astype('category')
loan_data["Loan_Status"] = loan_data["Loan_Status"].cat.codes
loan_data["Loan_Status"] = loan_data["Loan_Status"].astype('Int64')
loan_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                614 non-null   object
1   Gender                 614 non-null   object
2   Married                614 non-null   object
3   Dependents             614 non-null   Int64
4   Education              614 non-null   object
5   Self_Employed          614 non-null   object
6   ApplicantIncome         614 non-null   int64
7   CoapplicantIncome       614 non-null   float64
8   LoanAmount              614 non-null   float64
9   Loan_Amount_Term       614 non-null   Int64
10  Credit_History          614 non-null   object
11  Property_Area           614 non-null   object
12  Loan_Status             614 non-null   Int64
dtypes: Int64(3), float64(2), int64(1), object(7)
memory usage: 64.3+ KB

```



```
In [22]: loan_data_new = pd.concat([pd.get_dummies(loan_data[category_columns]),
loan_data[numerical_columns], loan_data["Loan_Status"]], axis=1)
for col_name in loan_data_new.columns:
    if(loan_data_new[col_name].dtype == 'object'):
        loan_data_new[col_name] = loan_data_new[col_name].astype('Int64')
loan_data_new
```

Out [22]:

	Gender_Female	Gender_Male	Gender_No Gender	Married_No	Married_Yes	Education_Graduate	Education_Not Graduate	Se
0	0	1	0	1	0	1	0	
1	0	1	0	0	1	1	0	
2	0	1	0	0	1	1	0	
3	0	1	0	0	1	0	1	
4	0	1	0	1	0	1	0	
...
609	1	0	0	1	0	1	0	
610	0	1	0	0	1	1	0	
611	0	1	0	0	1	1	0	
612	0	1	0	0	1	1	0	
613	1	0	0	1	0	1	0	

614 rows × 20 columns

```
In [23]: loan_data_new.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender_Female                        614 non-null    uint8
1   Gender_Male                         614 non-null    uint8
2   Gender_No Gender                    614 non-null    uint8
3   Married_No                         614 non-null    uint8
4   Married_Yes                        614 non-null    uint8
5   Education_Graduate                 614 non-null    uint8
6   Education_Not Graduate             614 non-null    uint8
7   Self_Employed_No                  614 non-null    uint8
8   Self_Employed_Yes                 614 non-null    uint8
9   Credit_History_0.0                 614 non-null    uint8
10  Credit_History_1.0                 614 non-null    uint8
11  Property_Area_Rural                 614 non-null    uint8
12  Property_Area_Semiurban             614 non-null    uint8
13  Property_Area_Urban                 614 non-null    uint8
14  Dependents                         614 non-null    Int64
15  LoanAmount                         614 non-null    float64
16  Loan_Amount_Term                   614 non-null    Int64
17  ApplicantIncome                    614 non-null    int64
18  CoapplicantIncome                  614 non-null    float64
19  Loan_Status                        614 non-null    Int64
dtypes: Int64(3), float64(2), int64(1), uint8(14)
memory usage: 39.1 KB
```

convert dataset to a feature matrix X

```
In [24]: X = loan_data_new.drop("Loan_Status", 1)
X.columns
```

```
Out[24]: Index(['Gender_Female', 'Gender_Male', 'Gender_No Gender', 'Married_No',
               'Married_Yes', 'Education_Graduate', 'Education_Not Graduate',
               'Self_Employed_No', 'Self_Employed_Yes', 'Credit_History_0.0',
               'Credit_History_1.0', 'Property_Area_Rural', 'Property_Area_Semiurban',
               'Property_Area_Urban', 'Dependents', 'LoanAmount', 'Loan_Amount_Term',
               'ApplicantIncome', 'CoapplicantIncome'],
              dtype='object')
```

```
In [25]: y = loan_data_new.iloc[:, -1].to_numpy()
```

normalizing feature matrix X

```
In [26]: def normalize(X):
          max_value = X.max()
          min_value = X.min()
          sub_value = max_value - min_value
          if(sub_value == 0):
              return 0
          return np.divide(np.subtract(X,min_value),sub_value)
```

```
In [27]: categorical_columns = list(set(loan_data_new.columns) -
set(numerical_columns) - set(["Loan_Status"]))
print(categorical_columns)

['Education_Graduate', 'Gender_Female', 'Credit_History_0.0', 'Property_Area_Semiurban',
'Married_Yes', 'Gender_No Gender', 'Self_Employed_No', 'Gender_Male', 'Self_Employed_Yes',
'Credit_History_1.0', 'Property_Area_Rural', 'Married_No', 'Property_Area_Urban',
'Education_Not Graduate']
```

```
In [28]: for col in numerical_columns:
          X[col] = normalize(X[col])
```

```
In [29]: X
```

```
Out[29]:
```

	Gender_Female	Gender_Male	Gender_No Gender	Married_No	Married_Yes	Education_Graduate	Education_Not Graduate	Se
0	0	1	0	1	0	1	0	
1	0	1	0	0	1	1	0	
2	0	1	0	0	1	1	0	
3	0	1	0	0	1	0	1	
4	0	1	0	1	0	1	0	
...	
609	1	0	0	1	0	1	0	
610	0	1	0	0	1	1	0	
611	0	1	0	0	1	1	0	
612	0	1	0	0	1	1	0	

	Gender_Female	Gender_Male	Gender_No Gender	Married_No	Married_Yes	Education_Graduate	Education_Not Graduate	Self_Employed
613	1	0	0	1	0	1	0	0

In [30]:

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender_Female                        614 non-null    uint8
1   Gender_Male                          614 non-null    uint8
2   Gender_No Gender                     614 non-null    uint8
3   Married_No                           614 non-null    uint8
4   Married_Yes                          614 non-null    uint8
5   Education_Graduate                   614 non-null    uint8
6   Education_Not Graduate                614 non-null    uint8
7   Self_Employed_No                     614 non-null    uint8
8   Self_Employed_Yes                    614 non-null    uint8
9   Credit_History_0.0                   614 non-null    uint8
10  Credit_History_1.0                   614 non-null    uint8
11  Property_Area_Rural                   614 non-null    uint8
12  Property_Area_Semiurban               614 non-null    uint8
13  Property_Area_Urban                   614 non-null    uint8
14  Dependents                           614 non-null    Float64
15  LoanAmount                           614 non-null    float64
16  Loan_Amount_Term                      614 non-null    Float64
17  ApplicantIncome                       614 non-null    float64
18  CoapplicantIncome                     614 non-null    float64
dtypes: Float64(2), float64(3), uint8(14)
memory usage: 33.7 KB
```

initial theta values

In [31]:

```
def transform_x_y(xx, yy):
    xx = np.c_[np.ones((xx.shape[0], 1)), xx]
    yy = yy[:, np.newaxis]
    m,n = np.shape(xx)
    one_vec = np.ones((m,1))
    xx = np.hstack((one_vec,xx))
    theta = np.zeros((n+1,1))
    return xx, yy, theta
```

In [32]:

```
old_X = X
old_Y = y
X, y, theta = transform_x_y(X, y)
```

gradient descent updation

```
In [33]: def sigmoid(z): # Activation function used to map any real value between 0
and 1
    z = z.astype(float)
    return (1 / (1 + np.exp(-z)))

def net_input(theta, x): # Computes the weighted sum of inputs
    return np.dot(x, theta)

def probability(theta, x): # Returns the probability after passing through
sigmoid
    return sigmoid(net_input(theta, x))

def cost_function(theta, x, y): # Computes the cost function for all the
training samples
    m = x.shape[0]
    total_cost = -(1 / m) * np.sum(y * np.log(probability(theta, x)) + (1 -
y) * np.log(1 - probability(theta, x)))
    return total_cost

def gradient(theta, x, y): # Computes the gradient of the cost function at
the point theta
    m = x.shape[0]
    return (1 / m) * np.dot(x.T, sigmoid(net_input(theta, x)) - y)

def gd_fit(x, y, theta):
    opt_weights = fmin_tnc(func=cost_function, x0=theta, fprime=gradient,
approx_grad=True, args=(x, y.flatten()))
    return opt_weights[0]
```

```
In [34]: parameters = gd_fit(X, y, theta)
parameters
```

```
Out[34]: array([ 0.34100623,  0.34100623,  0.04513421,  0.02795371, -0.19653694,
-0.36982231,  0.13315994, -0.09021046, -0.44293368,  0.17794204,
 0.10799482, -2.73525094,  1.22465379, -0.87600508,  0.02020988,
-0.68273099,  0.31420137, -0.41142908,  0.03601239,  0.02465119,
-1.31709396])
```

```
In [35]: def predict_gd(x):
    theta = parameters[:, np.newaxis]
    return probability(theta, x)

def accuracy_gd(x, actual_classes, probab_threshold=0.5):
    predicted_classes = (predict_gd(x) >= probab_threshold).astype(int)
    predicted_classes = predicted_classes.flatten()
    accuracy = np.mean(predicted_classes == actual_classes)
    return accuracy * 100, predicted_classes
```

```
In [36]: gradient_descent_accuracy, y_pred = accuracy_gd(X, y.flatten())
print("Accuracy using Gradient Descent = ", gradient_descent_accuracy, "%")
```

Accuracy using Gradient Descent = 81.27035830618892 %

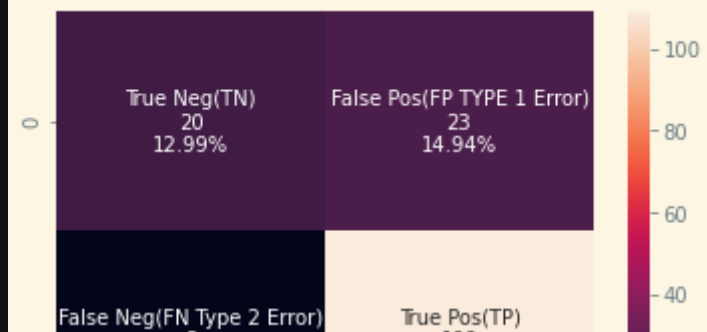
```
In [37]: ## Training model using only train dataset and then predicting test dataset
gd_x_train, gd_x_test, gd_y_train, gd_y_test = train_test_split(old_X,
old_Y, test_size=0.25, random_state=0)
gd_x_train, gd_y_train, theta = transform_x_y(gd_x_train, gd_y_train)
gd_parameters = gd_fit(gd_x_train, gd_y_train, theta)
print(gd_parameters)
gd_x_test, gd_y_test, theta = transform_x_y(gd_x_test, gd_y_test)
gradient_descent_accuracy_2, gd_y_pred = accuracy_gd(gd_x_test,
gd_y_test.flatten())
print("Accuracy using Gradient Descent after spliting dataset= ",
gradient_descent_accuracy_2, "%")
```

```
[ 0.46326678  0.46326678  0.1614313  0.24713049 -0.00123595 -0.1618432
 0.21285939  0.21346804 -0.34375405  0.31391866  0.3488988  -3.01717933
 0.79627855 -0.50348196  0.25558749 -0.3055587  0.2327919  -1.60039451
 -1.05512565  0.65278647 -0.73130014]
Accuracy using Gradient Descent after spliting dataset= 83.76623376623377 %
```

```
In [38]: gd_y_test = list(gd_y_test.flatten())
gd_y_pred = list(gd_y_pred)
print(classification_report(gd_y_test, gd_y_pred))
gd_cf_matrix = metrics.confusion_matrix(gd_y_test, gd_y_pred)
group_names = ['True Neg(TN)', 'False Pos(FP TYPE 1 Error)', 'False Neg(FN
Type 2 Error)', 'True Pos(TP)']
group_counts = ["{0:0.0f}".format(value) for value in
gd_cf_matrix.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
gd_cf_matrix.flatten()/np.sum(gd_cf_matrix)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
zip(group_names, group_counts, group_percentages)]
labels = np.asarray(labels).reshape(2,2)
import seaborn as sns
sns.heatmap(gd_cf_matrix, annot=labels, fmt='')
```

	precision	recall	f1-score	support
0	0.91	0.47	0.62	43
1	0.83	0.98	0.90	111
accuracy			0.84	154
macro avg	0.87	0.72	0.76	154
weighted avg	0.85	0.84	0.82	154

Out[38]: <AxesSubplot:>



7. Use sklearn built in function to find the model

```
In [39]: X = old_X
y = loan_data_new.iloc[:, -1]
y=y.astype('int')
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=0)
```

```
In [40]: logreg = LogisticRegression(max_iter = 10000,random_state=0)
logreg.fit(x_train, y_train)
y_pred = logreg.predict(x_test)
```

PART C: Performance Evaluation of the classifier

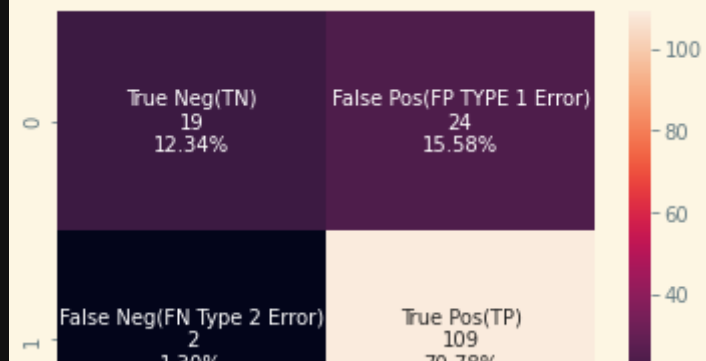
```
In [41]: parameters = logreg.coef_
print(parameters)

[[ 0.05805059  0.11849528 -0.17657126 -0.11954451  0.11951912  0.2417897
 -0.24181508  0.01133601 -0.0113614  -1.74853687  1.74851148 -0.31793508
  0.4163814  -0.09847171  0.26373488 -0.45327155 -0.34102166  0.24611225
 -0.87090286]]
```

8. Compute confusion matrix.

```
In [42]: cf_matrix = metrics.confusion_matrix(y_test, y_pred)
group_names = ['True Neg(TN)', 'False Pos (FP TYPE 1 Error)', 'False Neg (FN
Type 2 Error)', 'True Pos (TP)']
group_counts = ["{0:0.0f}".format(value) for value in cf_matrix.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
cf_matrix.flatten()/np.sum(cf_matrix)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
zip(group_names,group_counts,group_percentages)]
labels = np.asarray(labels).reshape(2,2)
import seaborn as sns
sns.heatmap(cf_matrix, annot=labels, fmt='')

Out[42]: <AxesSubplot:>
```



9. Compute the accuracy score.

```
In [43]: print("Accuracy using sklearn LogisticRegression()= ", logreg.score(x_test,
y_test)*100, "%")
```

Accuracy using sklearn LogisticRegression()= 83.11688311688312 %


10. Print a classification report using the following sklearn function

```
In [44]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.44	0.59	43
1	0.82	0.98	0.89	111
accuracy			0.83	154
macro avg	0.86	0.71	0.74	154
weighted avg	0.84	0.83	0.81	154

11. Plot ROC curve for loan status

```
In [45]: logit_roc_auc = roc_auc_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(x_test)[:,-1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' %
logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



12. Compare the performance of classifiers obtained in 6 and 7

In [46]:

```
print("Accuracy using Gradient Descent = ", gradient_descent_accuracy_2,
      "%")
print("Accuracy using sklearn LogisticRegression()= ", logreg.score(x_test,
y_test)*100, "%")
```

Accuracy using Gradient Descent = 83.76623376623377 %

Accuracy using sklearn LogisticRegression()= 83.11688311688312 %

Accuracy of Gradient Descent Logistic Regression is slightly better than
sklearn LogisticRegression()
Even precision, recall and f1-score are better in Gradient Descent Logistic
Regression than sklearn LogisticRegression()