

A Comprehensive Guide to Understand and Implement Text Classification in Python

[CLASSIFICATION](#)[DATA SCIENCE](#)[INTERMEDIATE](#)[MACHINE LEARNING](#)[NLP](#)[PROJECT](#)[PYTHON](#)[SUPERVISED](#)[TEXT](#)[UNSTRUCTURED DATA](#)

Introduction

One of the widely used [natural language processing](#) task in different business problems is “Text Classification”. The goal of text classification is to automatically classify the text documents into one or more defined categories. Some examples of text classification are:

- Understanding audience sentiment from social media,
- Detection of spam and non-spam emails,
- Auto tagging of customer queries, and
- Categorization of news articles into defined topics.

If you're a beginner in NLP, then you've come to the right place! We have designed a [comprehensive NLP course](#) just for you. It is one of our most popular courses and is just the right guide to kick start your NLP journey.

Project to apply Text Classification

Problem Statement


The objective of this task is to detect hate speech in tweets. For the sake of simplicity, we say a tweet contains hate speech if it has a racist or sexist sentiment associated with it. So, the task is to classify racist or sexist tweets from other tweets.

Formally, given a training sample of tweets and labels, where label '1' denotes the tweet is racist/sexist and label '0' denotes the tweet is not racist/sexist, your objective is to predict the labels on the test dataset.

[Practice Now](#)

Table of Contents

In this article, I will explain about the text classification and the step by step process to implement it in python.



Text Classification is an example of supervised machine learning task since a labelled dataset containing text documents and their labels is used for train a classifier. An end-to-end text classification pipeline is composed of three main components:

1. Dataset Preparation: The first step is the Dataset Preparation step which includes the process of loading a dataset and performing basic pre-processing. The dataset is then splitted into train and validation sets.

2. Feature Engineering: The next step is the Feature Engineering in which the raw dataset is transformed into flat features which can be used in a machine learning model. This step also includes the process of creating new features from the existing data.

3. Model Training: The final step is the Model Building step in which a machine learning model is trained on a labelled dataset.

4. Improve Performance of Text Classifier: In this article, we will also look at the different ways to improve the performance of text classifiers.

Note : *This article does not narrate NLP tasks in depth. If you want to revise the basics and come back here, you can always go through [this article](#).*

Getting your machine ready

Lets implement basic components in a step by step manner in order to create a text classification framework in python. To start with, import all the required libraries.

You would need requisite libraries to run this code – you can install them at their individual official links

- [Pandas](#)
- [Scikit-learn](#)
- [XGBoost](#)
- [TextBlob](#)
- [Keras](#)

```
# libraries for dataset preparation, feature engineering, model training
```

```
from sklearn import model_selection, preprocessing, linear_model, naive_bayes, metrics, svm
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn import decomposition, ensemble
import pandas, xgboost, numpy, textblob, string
from keras.preprocessing import text, sequence
from keras import layers, models, optimizers
```

1. Dataset preparation

For the purpose of this article, I am the using dataset of amazon reviews which can be [downloaded at this link](#). The dataset consists of 3.6M text reviews and their labels, we will use only a small fraction of data. To prepare the dataset, load the downloaded data into a pandas dataframe containing two columns – text and label. ([Source](#))

```
# load the dataset
data = open('data/corpus').read()
labels, texts = [], []
for i, line in enumerate(data.split("\n")):
    content = line.split()
    labels.append(content[0])
    texts.append(" ".join(content[1:]))
# create a dataframe using texts and lables
trainDF = pandas.DataFrame()
trainDF['text'] = texts
trainDF['label'] = labels
```

Next, we will split the dataset into training and validation sets so that we can train and test classifier. Also, we will encode our target column so that it can be used in machine learning models.

```
# split the dataset into training and validation datasets
train_x, valid_x, train_y, valid_y = model_selection.train_test_split(trainDF['text'], trainDF['label'])
# label encode the target variable
encoder = preprocessing.LabelEncoder()
train_y = encoder.fit_transform(train_y)
valid_y = encoder.fit_transform(valid_y)
```

2. Feature Engineering

The next step is the feature engineering step. In this step, raw text data will be transformed into feature vectors and new features will be created using the existing dataset. We will implement the following different ideas in order to obtain relevant features from our dataset.

2.1 Count Vectors as features

2.2 TF-IDF Vectors as features

- Word level
- N-Gram level
- Character level

2.3 Word Embeddings as features

2.4 Text / NLP based features

2.5 Topic Models as features

Lets look at the implementation of these ideas in detail.

2.1 Count Vectors as features

Count Vector is a matrix notation of the dataset in which every row represents a document from the corpus, every column represents a term from the corpus, and every cell represents the frequency count of a particular term in a particular document.

```
# create a count vectorizer object
count_vect = CountVectorizer(analyzer='word', token_pattern=r'\w{1,}')
count_vect.fit(trainDF['text'])
# transform the training and validation data using count vectorizer object
xtrain_count = count_vect.transform(train_x)
xvalid_count = count_vect.transform(valid_x)
```

2.2 TF-IDF Vectors as features

TF-IDF score represents the relative importance of a term in the document and the entire corpus. TF-IDF score is composed by two terms: the first computes the normalized Term Frequency (TF), the second term is the Inverse Document Frequency (IDF), computed as the logarithm of the number of the documents in the corpus divided by the number of documents where the specific term appears.

$TF(t) = (\text{Number of times term } t \text{ appears in a document}) / (\text{Total number of terms in the document})$

$IDF(t) = \log_e(\text{Total number of documents} / \text{Number of documents with term } t \text{ in it})$

TF-IDF Vectors can be generated at different levels of input tokens (words, characters, n-grams)

a. Word Level TF-IDF : Matrix representing tf-idf scores of every term in different documents

b. N-gram Level TF-IDF : N-grams are the combination of N terms together. This Matrix representing tf-idf scores of N-grams

c. Character Level TF-IDF : Matrix representing tf-idf scores of character level n-grams in the corpus

```
# word level tf-idf
tfidf_vect = TfidfVectorizer(analyzer='word', token_pattern=r'\w{1,}', max_features=5000)
tfidf_vect.fit(trainDF['text'])
xtrain_tfidf = tfidf_vect.transform(train_x)
xvalid_tfidf = tfidf_vect.transform(valid_x)

# ngram level tf-idf
tfidf_vect_ngram = TfidfVectorizer(analyzer='word', token_pattern=r'\w{1,}', ngram_range=(2,3), max_features=5000)
tfidf_vect_ngram.fit(trainDF['text'])
xtrain_tfidf_ngram = tfidf_vect_ngram.transform(train_x)
xvalid_tfidf_ngram = tfidf_vect_ngram.transform(valid_x)

# characters level tf-idf
tfidf_vect_ngram_chars = TfidfVectorizer(analyzer='char', token_pattern=r'\w{1,}', ngram_range=(2,3), max_features=5000)
tfidf_vect_ngram_chars.fit(trainDF['text'])
xtrain_tfidf_ngram_chars = tfidf_vect_ngram_chars.transform(train_x)
xvalid_tfidf_ngram_chars = tfidf_vect_ngram_chars.transform(valid_x)
```

2.3 Word Embeddings

A word embedding is a form of representing words and documents using a dense vector representation. The position of a word within the vector space is learned from text and is based on the words that surround the word when it is used. Word embeddings can be trained using the input corpus itself or can be generated using pre-trained word embeddings such as **Glove**, **FastText**, and **Word2Vec**. Any one of them can be downloaded and used as transfer learning. One can read more about word embeddings [here](#).

Following snippet shows how to use pre-trained word embeddings in the model. There are four essential steps:

1. Loading the pretrained word embeddings
2. Creating a tokenizer object
3. Transforming text documents to sequence of tokens and pad them
4. Create a mapping of token and their respective embeddings

You can download the pre-trained word embeddings from [here](#)

```
# load the pre-trained word-embedding vectors
embeddings_index = {}
for i, line in enumerate(open('data/wiki-news-300d-1M.vec')):
    values = line.split()
    embeddings_index[values[0]] = numpy.asarray(values[1:], dtype='float32')

# create a tokenizer
token = text.Tokenizer()
token.fit_on_texts(trainDF['text'])

word_index = token.word_index
# convert text to sequence of tokens and pad them to ensure equal length vectors
train_seq_x = sequence.pad_sequences(token.texts_to_sequences(train_x), maxlen=70)
valid_seq_x = sequence.pad_sequences(token.texts_to_sequences(valid_x), maxlen=70)

# create token-embedding mapping
embedding_matrix = numpy.zeros((len(word_index) + 1, 300))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
```

2.4 Text / NLP based features

A number of extra text based features can also be created which sometimes are helpful for improving text classification models. Some examples are:

1. Word Count of the documents – total number of words in the documents
2. Character Count of the documents – total number of characters in the documents
3. Average Word Density of the documents – average length of the words used in the documents
4. Punctuation Count in the Complete Essay – total number of punctuation marks in the documents
5. Upper Case Count in the Complete Essay – total number of upper count words in the documents
6. Title Word Count in the Complete Essay – total number of proper case (title) words in the documents
7. Frequency distribution of Part of Speech Tags:
 - Noun Count
 - Verb Count
 - Adjective Count
 - Adverb Count
 - Pronoun Count

These features are highly experimental ones and should be used according to the problem statement only.

```
trainDF['char_count'] = trainDF['text'].apply(len) trainDF['word_count'] = trainDF['text'].apply(lambda x: len(x.split())) trainDF['word_density'] = trainDF['char_count'] / (trainDF['word_count']+1) trainDF['punctuation_count'] = trainDF['text'].apply(lambda x: len("".join(_ for _ in x if _ in string.punctuation))) trainDF['title_word_count'] = trainDF['text'].apply(lambda x: len([wrd for wrd in x.split() if wrd.istitle()])) trainDF['upper_case_word_count'] = trainDF['text'].apply(lambda x: len([wrd for wrd in x.split() if wrd.isupper()])))
```

```
pos_family = { 'noun' : ['NN','NNS','NNP','NNPS'], 'pron' : ['PRP','PRP$','WP','WP$'], 'verb' : ['VB','VBD','VBG','VBN','VBP','VBZ'], 'adj' : ['JJ','JJR','JJS'], 'adv' : ['RB','RBR','RBS','WRB'] } # function to check and get the part of speech tag count of a words in a given sentence def check_pos_tag(x, flag): cnt = 0 try: wiki = textblob.TextBlob(x) for tup in wiki.tags: ppo = list(tup)[1] if ppo in pos_family[flag]: cnt += 1 except: pass return cnt trainDF['noun_count'] = trainDF['text'].apply(lambda x: check_pos_tag(x, 'noun')) trainDF['verb_count'] = trainDF['text'].apply(lambda x: check_pos_tag(x, 'verb')) trainDF['adj_count'] = trainDF['text'].apply(lambda x: check_pos_tag(x, 'adj')) trainDF['adv_count'] = trainDF['text'].apply(lambda x: check_pos_tag(x, 'adv')) trainDF['pron_count'] = trainDF['text'].apply(lambda x: check_pos_tag(x, 'pron'))
```

2.5 Topic Models as features

Topic Modelling is a technique to identify the groups of words (called a topic) from a collection of documents that contains best information in the collection. I have used **Latent Dirichlet Allocation** for generating Topic Modelling Features. LDA is an iterative model which starts from a fixed number of topics. Each topic is represented as a distribution over words, and each document is then represented as a distribution over topics. Although the tokens themselves are meaningless, the probability distributions over words provided by the topics provide a sense of the different ideas contained in the documents. One can read more about topic modelling [here](#)

Lets see its implementation:

```
# train a LDA Model lda_model = decomposition.LatentDirichletAllocation(n_components=20,
learning_method='online', max_iter=20) X_topics = lda_model.fit_transform(xtrain_count) topic_word =
lda_model.components_ vocab = count_vect.get_feature_names() # view the topic models n_top_words = 10
topic_summaries = [] for i, topic_dist in enumerate(topic_word): topic_words = numpy.array(vocab)
[numpy.argsort(topic_dist)][:-(n_top_words+1):-1] topic_summaries.append(' '.join(topic_words))
```

3. Model Building

The final step in the text classification framework is to train a classifier using the features created in the previous step. There are many different choices of machine learning models which can be used to train a final model. We will implement following different classifiers for this purpose:

1. Naive Bayes Classifier
2. Linear Classifier
3. Support Vector Machine
4. Bagging Models
5. Boosting Models
6. Shallow Neural Networks
7. Deep Neural Networks
 - Convolutional Neural Network (CNN)
 - Long Short Term Modelr (LSTM)
 - Gated Recurrent Unit (GRU)
 - Bidirectional RNN
 - Recurrent Convolutional Neural Network (RCNN)
 - Other Variants of Deep Neural Networks

Lets implement these models and understand their details. The following function is a utility function which can be used to train a model. It accepts the classifier, feature_vector of training data, labels of training data and feature vectors of valid data as inputs. Using these inputs, the model is trained and accuracy score is computed.

```
def train_model(classifier, feature_vector_train, label, feature_vector_valid, is_neural_net=False): # fit
the training dataset on the classifier classifier.fit(feature_vector_train, label) # predict the labels on
validation dataset predictions = classifier.predict(feature_vector_valid) if is_neural_net: predictions =
predictions.argmax(axis=-1) return metrics.accuracy_score(predictions, valid_y)
```

3.1 Naive Bayes

Implementing a naive bayes model using sklearn implementation with different features

Naive Bayes is a classification technique based on Bayes' Theorem with an assumption of independence among predictors. A Naive Bayes classifier assumes that the presence of a particular feature in a class is unrelated to the presence of any other feature [here](#).

```
# Naive Bayes on Count Vectors accuracy = train_model(naive_bayes.MultinomialNB(), xtrain_count, train_y,
xvalid_count) print "NB, Count Vectors: ", accuracy # Naive Bayes on Word Level TF IDF Vectors accuracy =
train_model(naive_bayes.MultinomialNB(), xtrain_tfidf, train_y, xvalid_tfidf) print "NB, WordLevel TF-IDF: ",
accuracy # Naive Bayes on Ngram Level TF IDF Vectors accuracy = train_model(naive_bayes.MultinomialNB(),
xtrain_tfidf_ngram, train_y, xvalid_tfidf_ngram) print "NB, N-Gram Vectors: ", accuracy # Naive Bayes on
Character Level TF IDF Vectors accuracy = train_model(naive_bayes.MultinomialNB(), xtrain_tfidf_ngram_chars,
train_y, xvalid_tfidf_ngram_chars) print "NB, CharLevel Vectors: ", accuracy
```

```
NB, Count Vectors: 0.7004 NB, WordLevel TF-IDF: 0.7024 NB, N-Gram Vectors: 0.5344 NB, CharLevel Vectors:
0.6872
```

3.2 Linear Classifier

Implementing a Linear Classifier (Logistic Regression)

Logistic regression measures the relationship between the categorical dependent variable and one or more independent variables by estimating probabilities using a logistic/sigmoid function. One can read more about logistic regression [here](#)

```
# Linear Classifier on Count Vectors accuracy = train_model(linear_model.LogisticRegression(), xtrain_count,
train_y, xvalid_count) print "LR, Count Vectors: ", accuracy # Linear Classifier on Word Level TF IDF Vectors
accuracy = train_model(linear_model.LogisticRegression(), xtrain_tfidf, train_y, xvalid_tfidf) print "LR,
WordLevel TF-IDF: ", accuracy # Linear Classifier on Ngram Level TF IDF Vectors accuracy =
train_model(linear_model.LogisticRegression(), xtrain_tfidf_ngram, train_y, xvalid_tfidf_ngram) print "LR, N-
Gram Vectors: ", accuracy # Linear Classifier on Character Level TF IDF Vectors accuracy =
train_model(linear_model.LogisticRegression(), xtrain_tfidf_ngram_chars, train_y, xvalid_tfidf_ngram_chars)
print "LR, CharLevel Vectors: ", accuracy
```

```
LR, Count Vectors: 0.7048 LR, WordLevel TF-IDF: 0.7056 LR, N-Gram Vectors: 0.4896 LR, CharLevel Vectors:
0.7012
```

3.3 Implementing a SVM Model

Support Vector Machine (SVM) is a supervised machine learning algorithm which can be used for both classification or regression challenges. The model extracts a best possible hyper-plane / line that segregates the two classes. One can read more about it [here](#)

```
# SVM on Ngram Level TF IDF Vectors accuracy = train_model(svm.SVC(), xtrain_tfidf_ngram, train_y,
xvalid_tfidf_ngram) print "SVM, N-Gram Vectors: ", accuracy
```

```
SVM, N-Gram Vectors: 0.5296
```

3.4 Bagging Model

Implementing a Random Forest Model

Random Forest models are a type of ensemble models, particularly bagging models. They are part of the tree based model family. One can read more about Bagging and random forests [here](#)

```
# RF on Count Vectors accuracy = train_model(ensemble.RandomForestClassifier(), xtrain_count, train_y,
xvalid_count) print "RF, Count Vectors: ", accuracy # RF on Word Level TF IDF Vectors accuracy =
train_model(ensemble.RandomForestClassifier(), xtrain_tfidf, train_y, xvalid_tfidf) print "RF, WordLevel TF-
IDF: ", accuracy
```

RF, Count Vectors: 0.6972 RF, WordLevel TF-IDF: 0.6988

3.5 Boosting Model

Implementing Xtereme Gradient Boosting Model

Boosting models are another type of ensemble models part of tree based models. Boosting is a machine learning ensemble meta-algorithm for primarily reducing bias, and also variance in supervised learning, and a family of machine learning algorithms that convert weak learners to strong ones. A weak learner is defined to be a classifier that is only slightly correlated with the true classification (it can label examples better than random guessing). Read more about these models [here](#)

```
# Extereme Gradient Boosting on Count Vectors accuracy = train_model(xgboost.XGBClassifier(),
xtrain_count.tocsc(), train_y, xvalid_count.tocsc()) print "Xgb, Count Vectors: ", accuracy # Extereme
Gradient Boosting on Word Level TF IDF Vectors accuracy = train_model(xgboost.XGBClassifier(),
xtrain_tfidf.tocsc(), train_y, xvalid_tfidf.tocsc()) print "Xgb, WordLevel TF-IDF: ", accuracy # Extereme
Gradient Boosting on Character Level TF IDF Vectors accuracy = train_model(xgboost.XGBClassifier(),
xtrain_tfidf_ngram_chars.tocsc(), train_y, xvalid_tfidf_ngram_chars.tocsc()) print "Xgb, CharLevel Vectors:
", accuracy
```


```
/usr/local/lib/python2.7/dist-packages/sklearn/preprocessing/label.py:151: DeprecationWarning: The truth
value of an empty array is ambiguous. Returning False, but in future this will result in an error. Use
`array.size > 0` to check that an array is not empty. if diff: /usr/local/lib/python2.7/dist-
packages/sklearn/preprocessing/label.py:151: DeprecationWarning: The truth value of an empty array is
ambiguous. Returning False, but in future this will result in an error. Use `array.size > 0` to check that an
array is not empty. if diff:
```

Xgb, Count Vectors: 0.6324 Xgb, WordLevel TF-IDF: 0.6364 Xgb, CharLevel Vectors: 0.6548

```
/usr/local/lib/python2.7/dist-packages/sklearn/preprocessing/label.py:151: DeprecationWarning: The truth
value of an empty array is ambiguous. Returning False, but in future this will result in an error. Use
`array.size > 0` to check that an array is not empty. if diff:
```

3.6 Shallow Neural Networks

A neural network is a mathematical model that is designed to behave similar to biological neurons and nervous system. These models are used to recognize complex patterns and relationships that exists within a labelled data. A shallow neural network contains mainly three types of layers – input layer, hidden layer, and output layer. Read more about neural networks [here](#)




```
def create_model_architecture(input_size): # create input layer
    input_layer = layers.Input((input_size, ), sparse=True) # create hidden layer
    hidden_layer = layers.Dense(100, activation="relu")(input_layer) # create output layer
    output_layer = layers.Dense(1, activation="sigmoid")(hidden_layer) classifier = models.Model(inputs = input_layer, outputs = output_layer)
    classifier.compile(optimizer=optimizers.Adam(), loss='binary_crossentropy')
    return classifier
classifier = create_model_architecture(xtrain_tfidf_ngram.shape[1])
accuracy = train_model(classifier, xtrain_tfidf_ngram, train_y, xvalid_tfidf_ngram, is_neural_net=True)
print "NN, Ngram Level TF IDF Vectors", accuracy
```

```
Epoch 1/1 7500/7500 [=====] - 1s 67us/step - loss: 0.6909 NN, Ngram Level TF IDF Vectors 0.5296
```


3.7 Deep Neural Networks

Deep Neural Networks are more complex neural networks in which the hidden layers performs much more complex operations than simple sigmoid or relu activations. Different types of deep learning models can be applied in text classification problems.



3.7.1 Convolutional Neural Network

In Convolutional neural networks, convolutions over the input layer are used to compute the output. This results in local connections, where each region of the input is connected to a neuron in the output. Each layer applies different filters and combines their results.



Read more about Convolutional Neural Networks [here](#)


```
def create_cnn(): # Add an Input Layer
    input_layer = layers.Input((70, )) # Add the word embedding Layer
    embedding_layer = layers.Embedding(len(word_index) + 1, 300, weights=[embedding_matrix], trainable=False)
    (input_layer) embedding_layer = layers.SpatialDropout1D(0.3)(embedding_layer) # Add the convolutional Layer
    conv_layer = layers.Convolution1D(100, 3, activation="relu")(embedding_layer) # Add the pooling Layer
    pooling_layer = layers.GlobalMaxPool1D()(conv_layer) # Add the output Layers
    output_layer1 = layers.Dense(50, activation="relu")(pooling_layer)
    output_layer1 = layers.Dropout(0.25)(output_layer1)
    output_layer2 = layers.Dense(1, activation="sigmoid")(output_layer1) # Compile the model
    model = models.Model(inputs=input_layer, outputs=output_layer2)
    model.compile(optimizer=optimizers.Adam(), loss='binary_crossentropy')
    return model
classifier = create_cnn()
accuracy = train_model(classifier, train_seq_x, train_y, valid_seq_x, is_neural_net=True)
print "CNN, Word Embeddings", accuracy
```

Epoch 1/1 7500/7500 [=====] - 12s 2ms/step - loss: 0.5847 CNN, Word Embeddings 0.5296

3.7.2 Recurrent Neural Network – LSTM

Unlike Feed-forward neural networks in which activation outputs are propagated only in one direction, the activation outputs from neurons propagate in both directions (from inputs to outputs and from outputs to inputs) in Recurrent Neural Networks. This creates loops in the neural network architecture which acts as a 'memory state' of the neurons. This state allows the neurons an ability to remember what have been learned so far.

The memory state in RNNs gives an advantage over traditional neural networks but a problem called Vanishing Gradient is associated with them. In this problem, while learning with a large number of layers, it becomes really hard for the network to learn and tune the parameters of the earlier layers. To address this problem, A new type of RNNs called LSTMs (Long Short Term Memory) Models have been developed.



Read more about LSTMs [here](#)

```
def create_rnn_lstm():
```

```
# Add an Input Layer input_layer = layers.Input((70, )) # Add the word embedding Layer embedding_layer =
layers.Embedding(len(word_index) + 1, 300, weights=[embedding_matrix], trainable=False)(input_layer)
embedding_layer = layers.SpatialDropout1D(0.3)(embedding_layer) # Add the LSTM Layer lstm_layer =
layers.LSTM(100)(embedding_layer) # Add the output Layers output_layer1 = layers.Dense(50, activation="relu")
(lstm_layer) output_layer1 = layers.Dropout(0.25)(output_layer1) output_layer2 = layers.Dense(1,
activation="sigmoid")(output_layer1) # Compile the model model = models.Model(inputs=input_layer,
outputs=output_layer2) model.compile(optimizer=optimizers.Adam(), loss='binary_crossentropy') return model
classifier = create_rnn_lstm() accuracy = train_model(classifier, train_seq_x, train_y, valid_seq_x,
is_neural_net=True) print "RNN-LSTM, Word Embeddings", accuracy
```

```
Epoch 1/1 7500/7500 [=====] - 22s 3ms/step - loss: 0.6899 RNN-LSTM, Word Embeddings
0.5124
```

3.7.3 Recurrent Neural Network – GRU

Gated Recurrent Units are another form of recurrent neural networks. Lets add a layer of GRU instead of LSTM in our network.

```
def create_rnn_gru():
```

```
# Add an Input Layer input_layer = layers.Input((70, )) # Add the word embedding Layer embedding_layer =
layers.Embedding(len(word_index) + 1, 300, weights=[embedding_matrix], trainable=False)(input_layer)
embedding_layer = layers.SpatialDropout1D(0.3)(embedding_layer) # Add the GRU Layer lstm_layer =
layers.GRU(100)(embedding_layer) # Add the output Layers output_layer1 = layers.Dense(50, activation="relu")
(lstm_layer) output_layer1 = layers.Dropout(0.25)(output_layer1) output_layer2 = layers.Dense(1,
activation="sigmoid")(output_layer1) # Compile the model model = models.Model(inputs=input_layer,
outputs=output_layer2) model.compile(optimizer=optimizers.Adam(), loss='binary_crossentropy') return model
classifier = create_rnn_gru() accuracy = train_model(classifier, train_seq_x, train_y, valid_seq_x,
is_neural_net=True) print "RNN-GRU, Word Embeddings", accuracy
```

```
Epoch 1/1 7500/7500 [=====] - 19s 3ms/step - loss: 0.6898 RNN-GRU, Word Embeddings
0.5124
```

3.7.4 Bidirectional RNN

RNN layers can be wrapped in Bidirectional layers as well. Lets wrap our GRU layer in bidirectional layer.

```
def create_bidirectional_rnn():
```

```
# Add an Input Layer input_layer = layers.Input((70, )) # Add the word embedding Layer embedding_layer =
layers.Embedding(len(word_index) + 1, 300, weights=[embedding_matrix], trainable=False)(input_layer)
embedding_layer = layers.SpatialDropout1D(0.3)(embedding_layer) # Add the LSTM Layer lstm_layer =
layers.Bidirectional(layers.GRU(100))(embedding_layer) # Add the output Layers output_layer1 =
layers.Dense(50, activation="relu")(lstm_layer) output_layer1 = layers.Dropout(0.25)(output_layer1)
output_layer2 = layers.Dense(1, activation="sigmoid")(output_layer1) # Compile the model model =
models.Model(inputs=input_layer, outputs=output_layer2) model.compile(optimizer=optimizers.Adam(),
```

```
loss='binary_crossentropy') return model classifier = create_bidirectional_rnn() accuracy =
train_model(classifier, train_seq_x, train_y, valid_seq_x, is_neural_net=True) print "RNN-Bidirectional, Word
Embeddings", accuracy
```

```
Epoch 1/1 7500/7500 [=====] - 32s 4ms/step - loss: 0.6889 RNN-Bidirectional, Word
Embeddings 0.5124
```

3.7.5 Recurrent Convolutional Neural Network

Once the essential architectures have been tried out, one can try different variants of these layers such as recurrent convolutional neural network. Another variants can be:

1. Hierarchical Attention Networks
2. Sequence to Sequence Models with Attention
3. Bidirectional Recurrent Convolutional Neural Networks
4. CNNs and RNNs with more number of layers

```
def create_rcnn():
```

```
# Add an Input Layer input_layer = layers.Input((70, )) # Add the word embedding Layer embedding_layer =
layers.Embedding(len(word_index) + 1, 300, weights=[embedding_matrix], trainable=False)(input_layer)
embedding_layer = layers.SpatialDropout1D(0.3)(embedding_layer) # Add the recurrent layer rnn_layer =
layers.Bidirectional(layers.GRU(50, return_sequences=True))(embedding_layer) # Add the convolutional Layer
conv_layer = layers.Convolution1D(100, 3, activation="relu")(embedding_layer) # Add the pooling Layer
pooling_layer = layers.GlobalMaxPool1D()(conv_layer) # Add the output Layers output_layer1 = layers.Dense(50,
activation="relu")(pooling_layer) output_layer1 = layers.Dropout(0.25)(output_layer1) output_layer2 =
layers.Dense(1, activation="sigmoid")(output_layer1) # Compile the model model =
models.Model(inputs=input_layer, outputs=output_layer2) model.compile(optimizer=optimizers.Adam(),
loss='binary_crossentropy') return model classifier = create_rcnn() accuracy = train_model(classifier,
train_seq_x, train_y, valid_seq_x, is_neural_net=True) print "CNN, Word Embeddings", accuracy
```

```
Epoch 1/1 7500/7500 [=====] - 11s 1ms/step - loss: 0.6902 CNN, Word Embeddings
0.5124
```

Improving Text Classification Models

While the above framework can be applied to a number of text classification problems, but to achieve a good accuracy some improvements can be done in the overall framework. For example, following are some tips to improve the performance of text classification models and this framework.

1. Text Cleaning : text cleaning can help to reduce the noise present in text data in the form of stopwords, punctuations marks, suffix variations etc. This [article](#) can help to understand how to implement text classification in detail.

2. Hstacking Text / NLP features with text feature vectors : In the feature engineering section, we generated a number of different feature vectors, combining them together can help to improve the accuracy of the classifier.

3. Hyperparameter Tuning in modelling : Tuning the parameters is an important step, a number of parameters such as tree length, leafs, network parameters etc can be fine tuned to get a best fit model.

4. Ensemble Models : Stacking different models and blending their outputs can help to further improve the results. Read more about ensemble models [here](#)

Projects

Now, its time to take the plunge and actually play with some other real datasets. So are you ready to take on the challenge? Accelerate your NLP journey with the following Practice Problems:

	Practice Problem: Identify the Sentiments	Identify the sentiment of tweets
	Practice Problem : Twitter Sentiment Analysis	To detect hate speech in tweets


End Notes

In this article, we discussed about how to prepare a text dataset like cleaning/creating training and validation dataset, perform different types of feature engineering like Count Vector/TF-IDF/ Word Embedding/ Topic Modelling and basic text features, and finally trained a variety of classifiers like Naive Bayes/ Logistic regression/ SVM/ MLP/ LSTM and GRU. At the end, discussed about different approach to improve the performance of text classifiers.

Note: There is a video course, [Natural Language Processing using Python](#), with 3 real life projects, two of them involve text classification.

Did you find this article useful ? Share your views and opinions in the comments section below.

[Learn, compete, hack and get hired!](#)



Shivam Bansal

Shivam Bansal is a data scientist with exhaustive experience in Natural Language Processing and Machine Learning in several domains. He is passionate about learning and always looks forward to solving challenging analytical problems.